

# Exercise\_round\_1

February 4, 2025

## 1 Exercise 3. (Kalman Filtering example)

Python imports. We will not need more for this exercise. If these are not installed (for instance on Google Colab), simply run `!pip install <package>` into a code cell.

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import scipy.linalg as linalg

np.random.seed(0) # fix the seed so that results are the same for everyone
```

### 1.0.1 The filtering utilities

```
[2]: def kf_predict(m, P, A, Q, b=None):
    """Kalman filter prediction function.

    This computes the mean and covariance of the predictive distribution

    .. math::
        p(x_t \mid y_{1:t-1}) = \int_{x_{t-1}} p(x_t \mid x_{t-1}) p(x_{t-1} \mid
    \hookrightarrow y_{1:t-1}) dx_{t-1}

    in the case when :math: p(x_{t-1} \mid y_{1:t-1}) = N(x_{t-1}; m, P) is
    \hookrightarrow Gaussian and the transition model

    .. math::
        p(x_t \mid x_{t-1}) = N(x_t; A x_{t-1} + b, Q)

    is conditionally Gaussian.

    Parameters
    -----
    m: ndarray
        The prior mean.
    P: ndarray
        The prior covariance.
    A: ndarray
```

```

    The transition matrix.
Q: ndarray
    The transition noise covariance.
b: ndarray, optional
    An offset, if None is passed, it is assumed to be 0.

Returns
-----
m: ndarray
    The predictive mean.
P: ndarray
    The predictive covariance.
"""

m = A @ m
if b is not None:
    m = m + b

P = Q + A @ P @ A.T
return m, P

def kf_update(m, P, y, H, R, c=None):
    r"""Kalman filter update function.

    This computes the mean and covariance of the posterior distribution

    .. math::
        p(x_t \mid y_{1:t}) \propto p(y_t \mid x_t) p(x_t \mid y_{1:t-1})

    in the case when :math: p(x_t \mid y_{1:t-1}) = N(x_t; m, P) is
    Gaussian and the observation model

    .. math::
        p(y_t \mid x_t) = N(y_t; H x_t + c, R)

    is conditionally Gaussian.

    Parameters
    -----
    m: ndarray
        The prior mean.
    P: ndarray
        The prior covariance.
    H: ndarray
        The observation/emission matrix.
    R: ndarray

```

```

        The observation error covariance.
c: ndarray, optional
    An offset, if None is passed, it is assumed to be 0.

Returns
-----
m: ndarray
    The filtered mean.
P: ndarray
    The filtered covariance.
"""

# predictive mean of the observation
y_predicted = H @ m
if c is not None:
    y_predicted = y_predicted + c

# predictive error
residual = y - y_predicted

# predictive covariance of the observation
S = H @ P @ H.T + R

# Kalman gain
# This code  $P H^T / S$  by using the fact  $S$  is "symmetric positive definite",
↳ and solving instead of inverting.
K = linalg.solve(S, H @ P, assume_a='pos').T

# update
m = m + K @ residual
P = P - K @ S @ K.T

return m, P

def rts_update(m, P, m_f, P_f, A, Q, b=None):
    r"""Rauch-Tung-Striebel update function.

    This computes the mean and covariance of the posterior distribution

    .. math::
        p(x_t \mid y_{1:T})

    Parameters
    -----
    m: ndarray
        The smoothed mean of the next time step.

```

```

P: ndarray
    The smoothed covariance of the next time step.
m_f: ndarray
    The filtered mean of the current time step.
P_f: ndarray
    The filtered covariance of the current time step.
A: ndarray
    The transition matrix.
Q: ndarray
    The transition noise covariance.
b: ndarray, optional
    An offset, if None is passed, it is assumed to be 0.

Returns
-----
m: ndarray
    The smoothed mean.
P: ndarray
    The smoothed covariance.
"""

# predictive covariance
S = A @ P_f @ A.T + Q

# predictive error
predictive_mean = A @ m_f
if b is not None:
    predictive_mean = predictive_mean + b
residual = m - predictive_mean

# Kalman gain
# This code  $P H^T / S$  by using the fact  $S$  is "symmetric positive definite",
↳ and solving instead of inverting.
K = linalg.solve(S, A @ P_f, assume_a='pos').T

# update
m = m_f + K @ residual
P = P_f - K @ S @ K.T

return m, P

```

### 1.0.2 The discretisation function

```

[3]: def lti_discretization(F, Q, dt):
    r"""Discretises a linear time invariant stochastic differential equation
    ↳ using matrix exponentials.

```

```

Parameters
-----
F: ndarray
    The continuous transition matrix
Q: ndarray
    The covariance of the noise process
dt: float
    The time step considered

Returns
-----
A: ndarray
    The discretised transition matrix
Q: ndarray
    The discretised noise covariance matrix

"""
dim = F.shape[0]
zeros = np.zeros((dim, dim))
eye = np.eye(dim)

A = linalg.expm(F * dt)
Phi = np.block([[F, Q],
                [zeros, -F.T]])

AB = linalg.expm(Phi * dt) @ np.block([[zeros],
                                       [eye]])
Q = linalg.solve(AB[dim:].T, AB[:dim].T).T # AB[:dim] / AB[dim:]

return A, Q

```

## 1.1 (a) Sinusoidal example

We will be trying to “denoisify” a corrupted sine function. Consider a model given by  $x(t) = \sin(t)$  for all  $t$ , and  $y(t_k) = x(t_k) + \epsilon_{t_k}$  with i.i.d.  $\epsilon_{t_k} \sim N(0, \sigma^2)$  and  $t_0 = 0, t_k = t_{k+1} + \Delta$ .

```

[4]: sigma = 0.1
    delta = 0.1
    ts = np.arange(0, 30, delta) # [0, 0.1, 0.2, ..., 29.9]
    T = ts.shape[0]

    sines = np.sin(ts)
    sin_obs = sines + sigma * np.random.randn(T)

```

To recover the noisy signal, we can form a Wiener velocity type of state-space model containing the signal and its derivative as the states:

$$\frac{d^2x(t)}{dt^2} = w(t)$$

where  $w(t)$  is a white noise process.

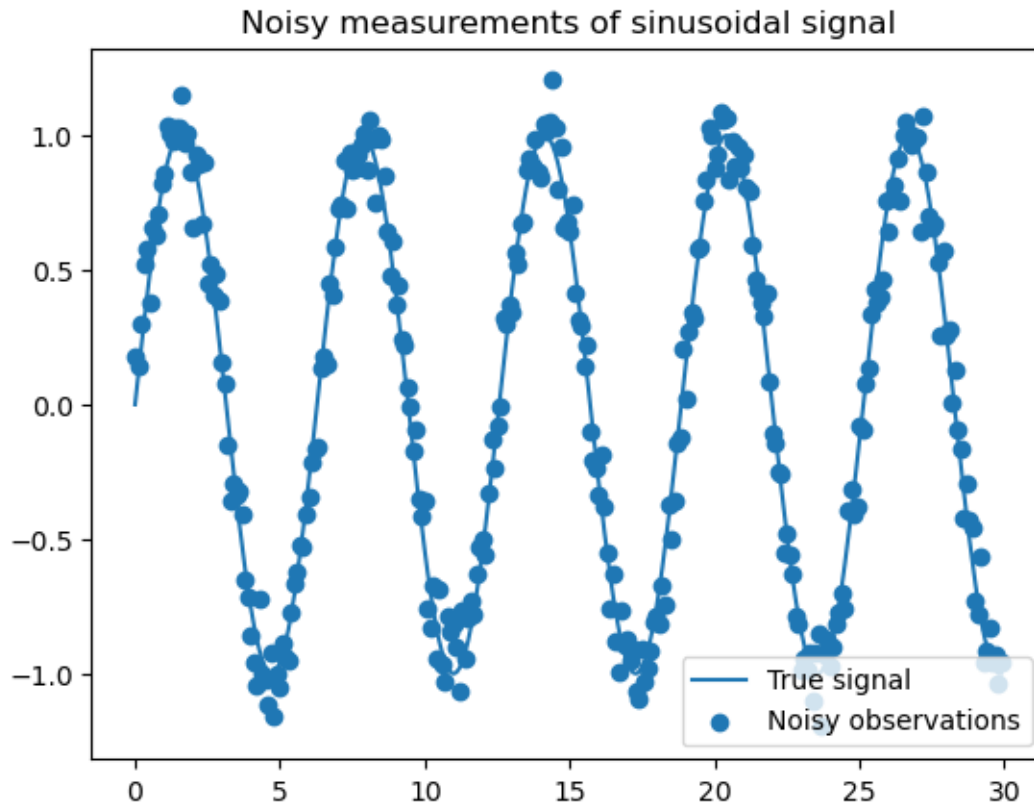
```
[5]: # transition model
F = np.array([[0., 1.],
              [0., 0.]])
q = 0.1
A, Q = lti_discretization(F, np.diag([0, q]), delta)

# observation model
H = np.array([[1., 0.]])
R = np.array([[sigma ** 2]])

# initial mean and covariance
m = np.array([0., 0.])
P = np.diag([0.1, 2])
```

```
[6]: plt.title("Noisy measurements of sinusoidal signal")
plt.plot(ts, sines, label="True signal")
plt.scatter(ts, sin_obs, label="Noisy observations")
plt.legend()
```

```
[6]: <matplotlib.legend.Legend at 0x7f0ec77a1660>
```



Let us now run a Kalman filter on the data

```
[7]: # Initialise the arrays to store the results
filtered_means = np.empty((T, 2))
filtered_covs = np.empty((T, 2, 2))

predicted_means = np.empty((T, 2))
predicted_covs = np.empty((T, 2, 2))

for i in range(T):
    t = ts[i]
    y = sin_obs[i]

    # Kalman filter prediction
    m, P = kf_predict(m, P, A, Q)
    predicted_means[i] = m
    predicted_covs[i] = P

    # Kalman filter update
    m, P = kf_update(m, P, y, H, R)
    filtered_means[i] = m
```

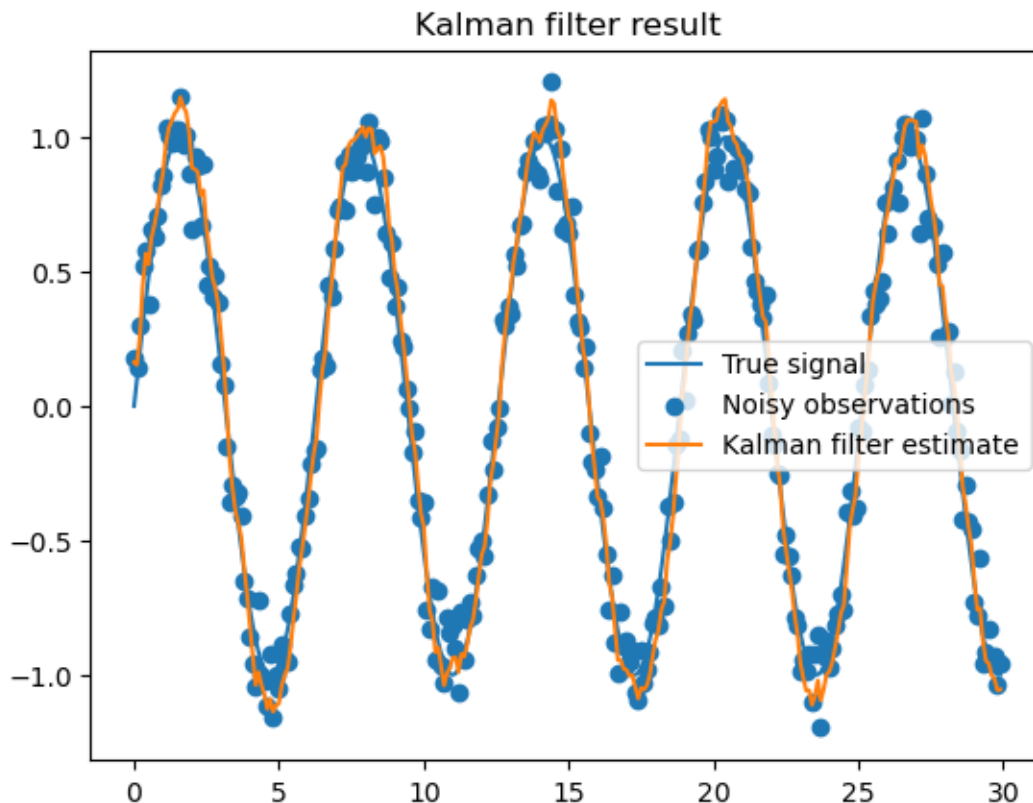
```

        filtered_covs[i] = P

plt.title("Kalman filter result")
plt.plot(ts, sines, label="True signal")
plt.scatter(ts, sin_obs, label="Noisy observations")
plt.plot(ts, filtered_means[:,0], label="Kalman filter estimate")
plt.legend()

```

[7]: <matplotlib.legend.Legend at 0x7f0eb876b400>



```

[8]: print(f"The RMSE of the filtering solution compared to the true data is {np.
      ↪mean((filtered_means[:, 0] - sines) ** 2) ** 0.5:.3f}")

```

The RMSE of the filtering solution compared to the true data is 0.082

```

[9]: print(f"Simply using the observations would result in {np.mean((sin_obs -
      ↪sines) ** 2) ** 0.5:.3f}")

```

Simply using the observations would result in 0.100

Run Rauch-Tung-Striebel smoother on the data.



```
[10]: # Initialise the arrays to store the results
smoothed_means = np.empty((T, 2))
smoothed_covs = np.empty((T, 2, 2))

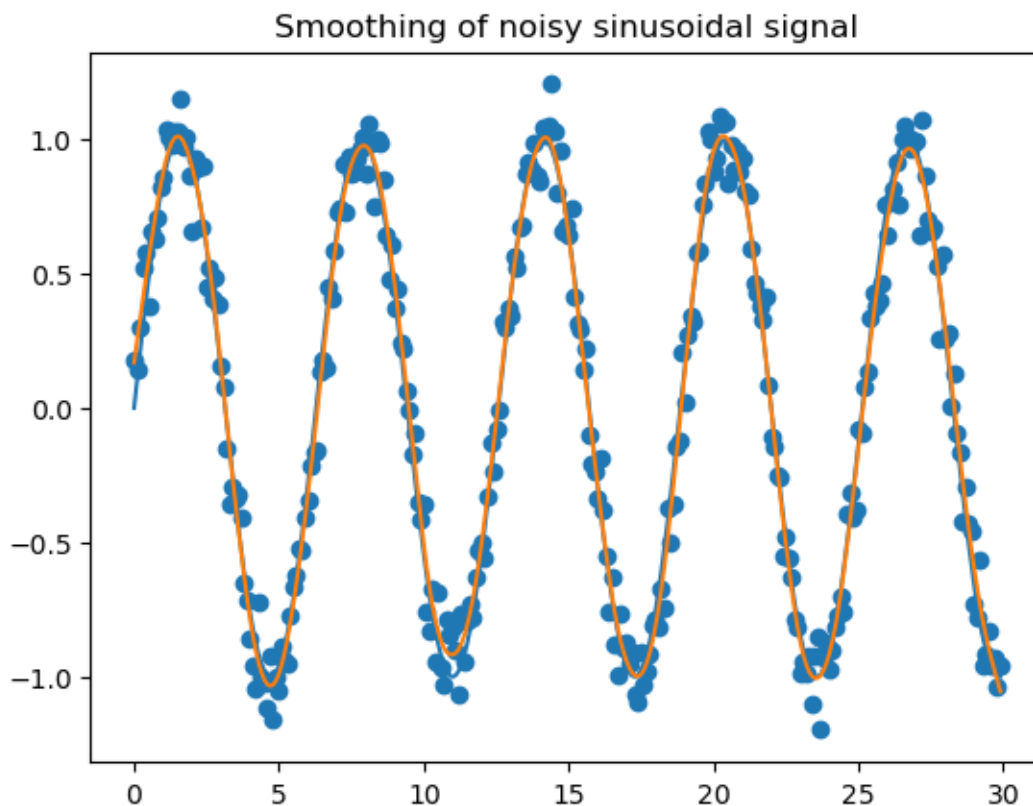
m = smoothed_means[-1] = filtered_means[-1]
P = smoothed_covs[-1] = filtered_covs[-1]

for i in range(T-2, -1, -1):
    t = ts[i]
    m_filtered = filtered_means[i]
    P_filtered = filtered_covs[i]

    m, P = rts_update(m, P, m_filtered, P_filtered, A, Q)
    smoothed_means[i] = m
    smoothed_covs[i] = P
```

```
[11]: plt.title("Smoothing of noisy sinusoidal signal")
plt.plot(ts, sines, label="True signal")
plt.scatter(ts, sin_obs, label="Noisy observations")
plt.plot(ts, smoothed_means[:,0], label="RTS smoother estimate")
```

```
[11]: [<matplotlib.lines.Line2D at 0x7f0eb8670340>]
```



```
[12]: print(f"The RMSE of the smoothing solution compared to the true data is {np.
      ↪mean((smoothed_means[:, 0] - sines) ** 2) ** 0.5:.3f}")
```

The RMSE of the smoothing solution compared to the true data is 0.037

## 1.2 Your turn.

### 1.2.1 (a)

Consider the following state space model:

$$\vec{x}_k = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \vec{x}_{k-1} + \vec{w}_{k-1} \quad (1)$$

$$y_k = \begin{pmatrix} 1 & 0 \end{pmatrix} \vec{x}_k + v_k \quad (2)$$

where  $\vec{x}_k = (x_k \dot{x}_k)^\top$  is the state,  $y_k$  is the measurement, and  $\vec{w}_k \sim \mathcal{N}(\vec{0}, \text{diag}(1/10^2, 1^2))$  and  $v_k \sim \mathcal{N}(0, 10^2)$  are white Gaussian noise processes.

Simulate a 100 step state sequence from the model and plot the signal  $x_k$ , signal derivative  $\dot{x}_k$  and the simulated measurements  $y_k$ . Start from an initial state drawn from a zero-mean 2d-Gaussian distribution with identity covariance.

```
[13]: import numpy as np
import matplotlib.pyplot as plt

# Define model parameters
A = np.array([[1., 1.], [0., 1.]]) # State transition matrix
H = np.array([[1., 0.]]) # Measurement model matrix

# State transition noise
w_mean = np.array([0., 0.])
w_var = np.diag([1/(10**2), 1**2])
w = np.array([0., 0.])

# Measurement noise
v_mean = 0
v_var = np.array([[10**2]])
v = 0

# SIMULATION

np.random.seed(0)

# Initial state
x0 = np.random.multivariate_normal(mean=np.array([0., 0.]), cov=np.array([[1., 0.], [0., 1.])))

# Number of time steps
```

```

n_steps = 100

# State vector
x = np.zeros((n_steps, 2))
x[0] = x0

# Measurement vector
y = np.zeros(n_steps)

for k in range(1, n_steps):

    w = np.random.multivariate_normal(mean=w_mean, cov=w_var)
    v = np.random.normal(loc=v_mean, scale=v_var)

    x[k] = A @ x[k-1] + w
    y[k] = H @ x[k] + v

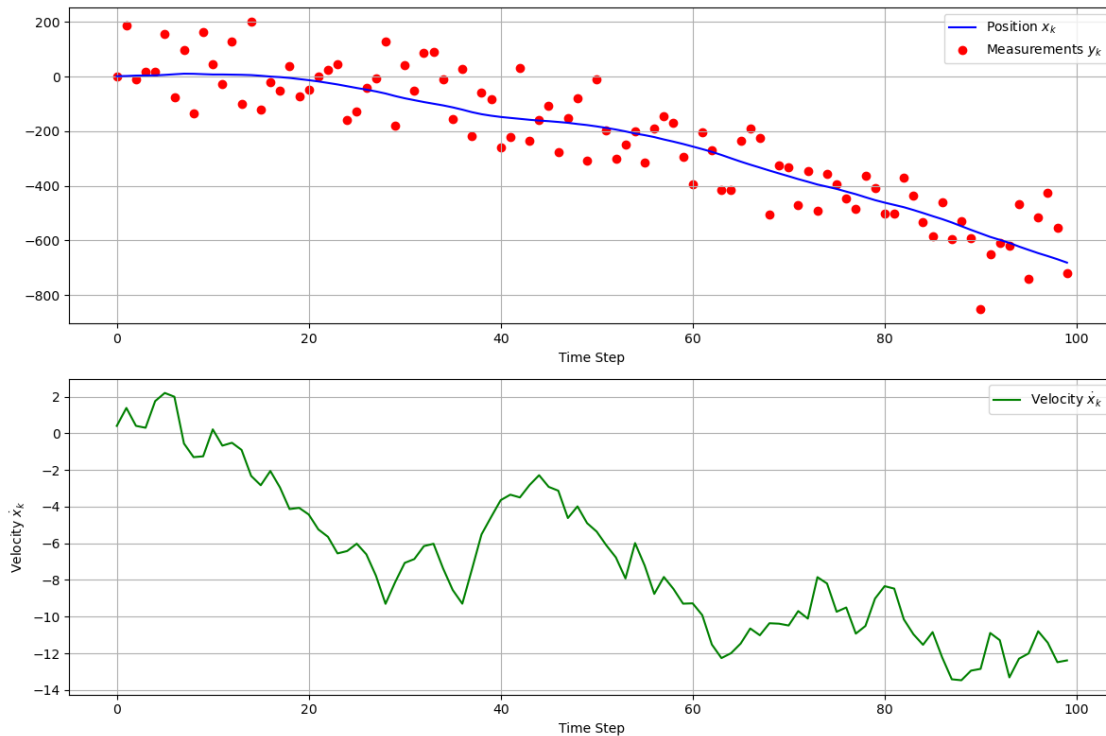
# Plotting
plt.figure(figsize=(12, 8))

# Plot position  $x_k$  and noisy measurements  $y_k$ 
plt.subplot(2, 1, 1)
plt.plot(range(n_steps), x[:, 0], label='Position  $x_k$ ', color='blue')
plt.scatter(range(n_steps), y, label='Measurements  $y_k$ ', color='red')
plt.xlabel('Time Step')
plt.grid(True)
plt.legend()

# Plot velocity  $\dot{x}_k$ 
plt.subplot(2, 1, 2)
plt.plot(range(n_steps), x[:, 1], label='Velocity  $\dot{x}_k$ ', color='green')
plt.xlabel('Time Step')
plt.ylabel('Velocity  $\dot{x}_k$ ')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

```



### 1.2.2 (b)

Use the Kalman filter for computing the state estimates.

Hint: This corresponds to writing a loop around the two functions

```
# initialize the mean and the covariance
m = ...
P = ...
for ...
    m, P = kf_predict(...)
    m, P = kf_update(...)
    # Store the result
```

already defined in a cell above on the output of the data generating process you have coded in (a).

```
[14]: m = np.array([0., 0.])
      P = np.diag([0.1, 1.])

      # Notating process and measurement noises for better context
      Q = w_var
      R = v_var

      predicted_means = np.empty((n_steps, 2))
      predicted_covs = np.empty((n_steps, 2, 2))
```

```

filtered_means = np.empty((n_steps, 2))
filtered_covs = np.empty((n_steps, 2, 2))

for k in range(1, n_steps):

    # Prediction step
    m, P = kf_predict(m, P, A, Q)
    predicted_means[k] = m
    predicted_covs[k] = P

    # Update step
    m, P = kf_update(m, P, y[k], H, R)
    filtered_means[k] = m
    filtered_covs[k] = P

```

### 1.2.3 (c)

Plot the state estimates  $\vec{m}_k$ , the true states  $\vec{x}_k$  and measurements  $y_k$  using [Matplotlib](#)

```

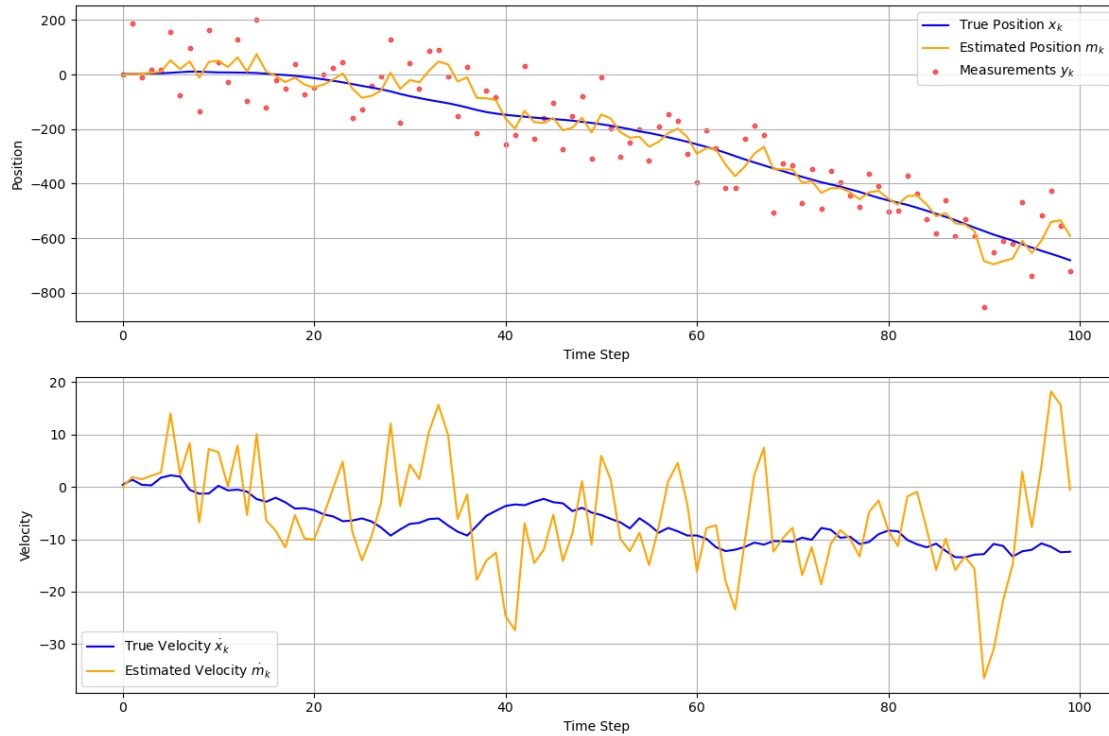
[15]: # Plotting results
fig, ax = plt.subplots(2, 1, figsize=(12, 8))

# Plot true position, estimated position, and measurements
ax[0].plot(range(n_steps), x[:, 0], label='True Position  $x_k$ ', color='blue')
ax[0].plot(range(n_steps), filtered_means[:, 0], label='Estimated Position  $\hat{x}_k$ ', color='orange')
ax[0].scatter(range(n_steps), y, label='Measurements  $y_k$ ', color='red', s=10, alpha=0.6)
ax[0].set_xlabel('Time Step')
ax[0].set_ylabel('Position')
ax[0].legend()
ax[0].grid(True)

# Plot true velocity and estimated velocity
ax[1].plot(range(n_steps), x[:, 1], label='True Velocity  $\dot{x}_k$ ', color='blue')
ax[1].plot(range(n_steps), filtered_means[:, 1], label='Estimated Velocity  $\dot{\hat{x}}_k$ ', color='orange')
ax[1].set_xlabel('Time Step')
ax[1].set_ylabel('Velocity')
ax[1].legend()
ax[1].grid(True)

plt.tight_layout()
plt.show()

```



#### 1.2.4 (d)

Compute the RMSE (root mean square error) of using the first components of vectors  $\vec{m}_k$  as the estimates of first components of states  $\vec{x}_k$ . Also compute the RMSE error that we would have if we used the measurements as the estimates.

```
[16]: # RMSE between Kalman estimates and true position
rmse_kalman = np.sqrt(np.mean((x[:, 0] - filtered_means[:, 0]) ** 2))

print(f"RMSE of Kalman estimates: {rmse_kalman:.4f}")

# RMSE between measurements and true position
rmse_measurements = np.sqrt(np.mean((x[:, 0] - y) ** 2))

print(f"RMSE of measurements: {rmse_measurements:.4f}")
```

```
RMSE of Kalman estimates: 48.1465
RMSE of measurements: 98.6949
```